

Data processing with the RTS

A GPU-accelerated calibration & imaging stream processor

Daniel Mitchell

2018 ICRAR/CASS Radio School

CSIRO ASTRONOMY AND SPACE SCIENCE
www.csiro.au



The RTS (Real-Time System)

A GPU-accelerated calibration & imaging stream processor designed for MWA

- Design Drivers
- Design Solutions
- Data & Processing Flow
- Using the RTS

Design Drivers

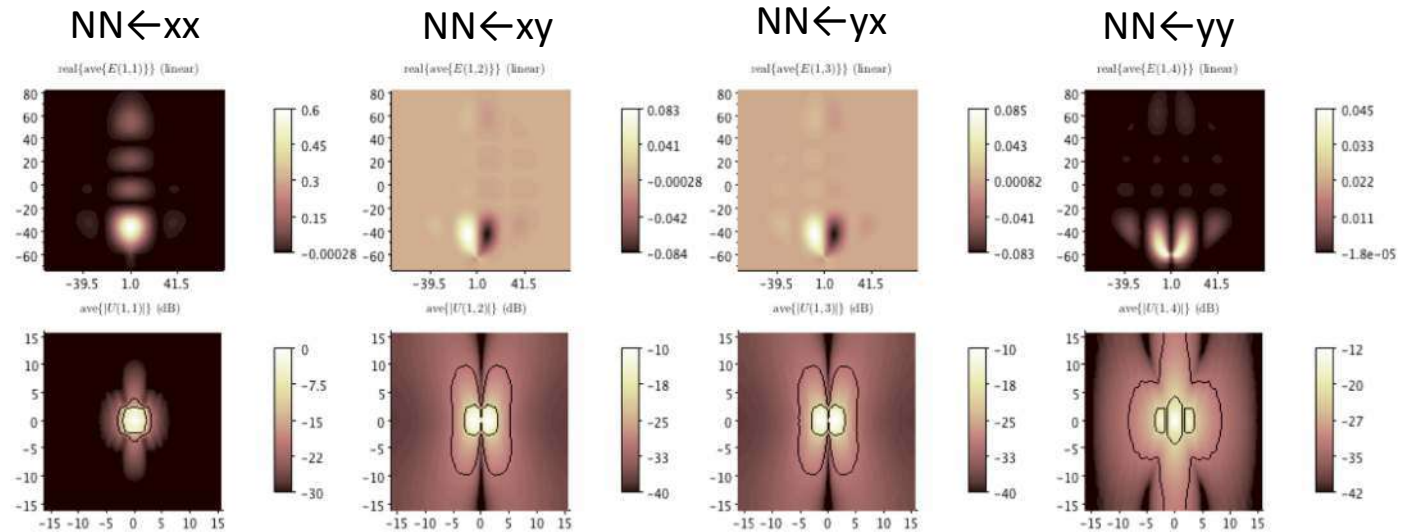
Design Drivers

- Very wide MWA field of view
 - Standard W-projection approach requires very large gridding kernels.
- Variable, highly polarised primary beams
 - Varying across the field of view and in time as a field is tracked.
 - Different tile beams due to analogue beamformer variability.
- Ionospheric refraction
 - Propagation through the ionosphere causes $\approx \lambda^2$ -dependent delays (λ -dependent phases). Kolmogorov turbulence + wave-like structures.
 - time-variable phase fluctuations across the aperture and the FoV.
- Challenging to couple very large W terms with highly variable A and I terms.
- Excellent MWA snapshot uv coverage and synthesised beam


Variable Polarised Primary Beams

All-sky polarised
tile response


Fourier
response



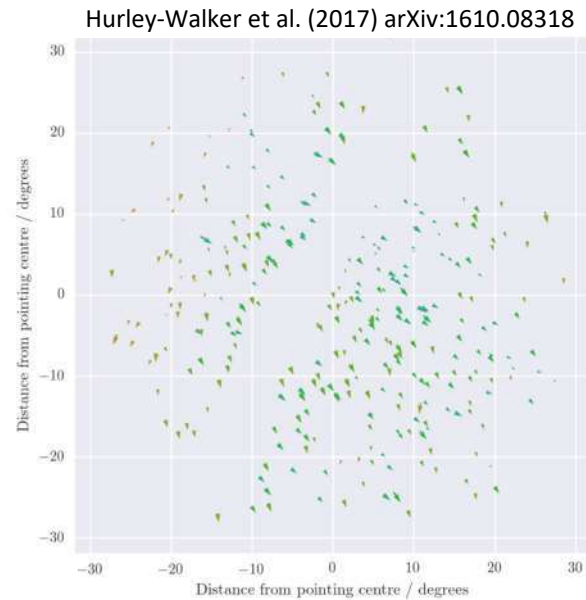


Time-dependent
tile response



Time-dependent
error



Ionospheric Refraction



Design Solutions


Accumulate Residual Snapshot Images

- Take advantage of MWA strengths: it is an ionospheric machine!
 - Lots of relatively short baselines.
 - Solve for the ionosphere rather than direction-dependent tile phases
 - increase SNR and decrease the number of free parameters.
 - Segment visibility data in time during imaging
 - ≈ linear ionospheric variations.
 - ≈ 2D phases in Fourier transform.
 - ≈ equal primary beams for each visibility.
- Well suited to stream processing and real-time operation
 - Parallelize in frequency for high-throughput stream processing.
 - Use a cluster of GPUs to reduce power and cost per FLOP.
 - Extend data reduction averaging times by moving to the image domain.

Peeling

- Subtract initial sky model
- Add strong sources back one-by-one, redo calibration for each and re-subtract

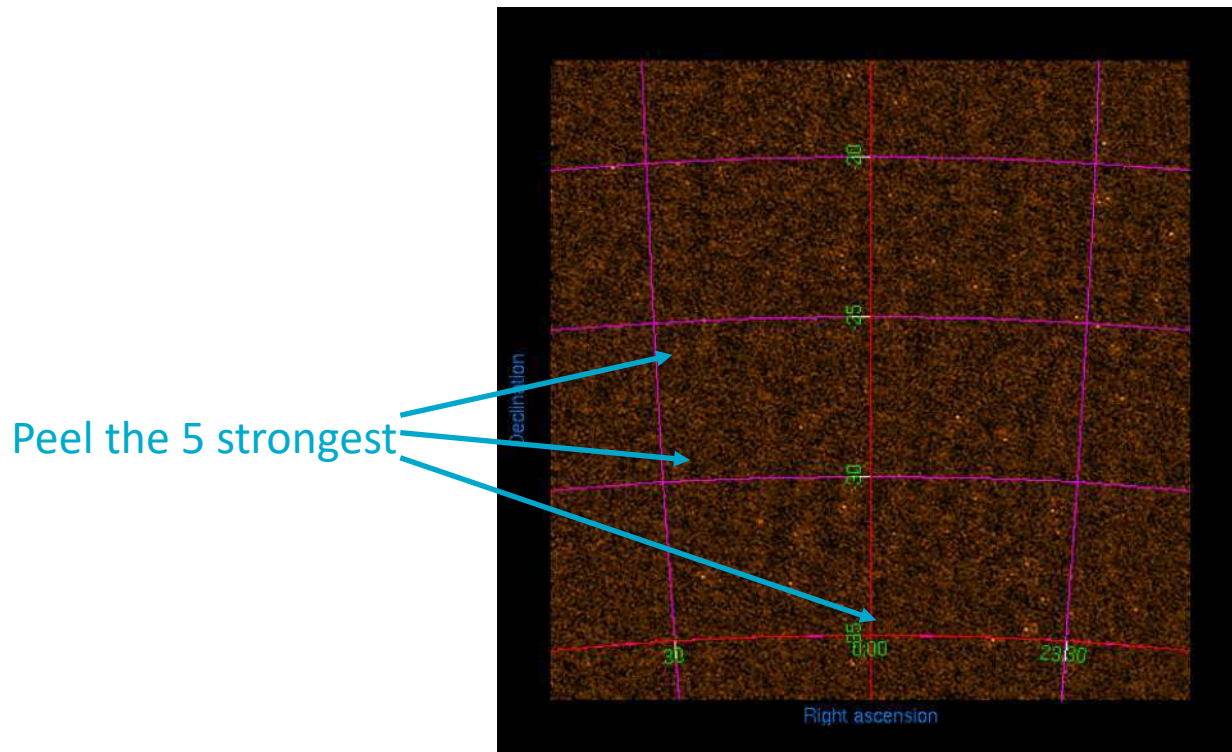
$$E_j = \left(\sum_{k \neq j} V_{jk} M_{jk}^\dagger \right) \left(\sum_{k \neq j} M_{jk} M_{jk}^\dagger \right)^{-1}; \quad M_{jk} = J_j S J_k^\dagger; \quad V_{jk} = M_{jk} + N_{jk}$$




Peeling

- Subtract initial sky model
- Add strong sources back one-by-one, redo calibration for each and re-subtract

$$E_j = (\sum_{k \neq j} V_{jk} M_{jk}^\dagger) (\sum_{k \neq j} M_{jk} M_{jk}^\dagger)^{-1}; \quad M_{jk} = J_j S J_k^\dagger; \quad V_{jk} = M_{jk} + N_{jk}$$




Peeling — over-peeling



Dirty image with direction-independent calibration


Peeling — over-peeling



Dirty image with direction-independent calibration and 50 sources subtracted (but not peeled)

The residuals are dominated by weaker sources, not by subtraction artefacts.


Peeling — over-peeling



Dirty image with direction-independent calibration and 50 sources subtracted (5 of the 50 peeled)

Peeling the 5 brightest sources doesn't have too much of an effect on the residuals.


Peeling — over-peeling



Dirty image with direction-independent calibration and 50 sources subtracted (10 of the 50 peeled)

But as more of the 50 sources are peeled the weaker sources and their sidelobes disappear!


Peeling — over-peeling



Dirty image with direction-independent calibration and 50 sources subtracted (20 of the 50 peeled)

But as more of the 50 sources are peeled the weaker sources and their sidelobes disappear!


Peeling — over-peeling



Dirty image with direction-independent calibration and 50 sources subtracted (30 of the 50 peeled)

But as more of the 50 sources are peeled the weaker sources and their sidelobes disappear!


Peeling — over-peeling



Dirty image with direction-independent calibration and 50 sources subtracted (40 of the 50 peeled)

But as more of the 50 sources are peeled the weaker sources and their sidelobes disappear!

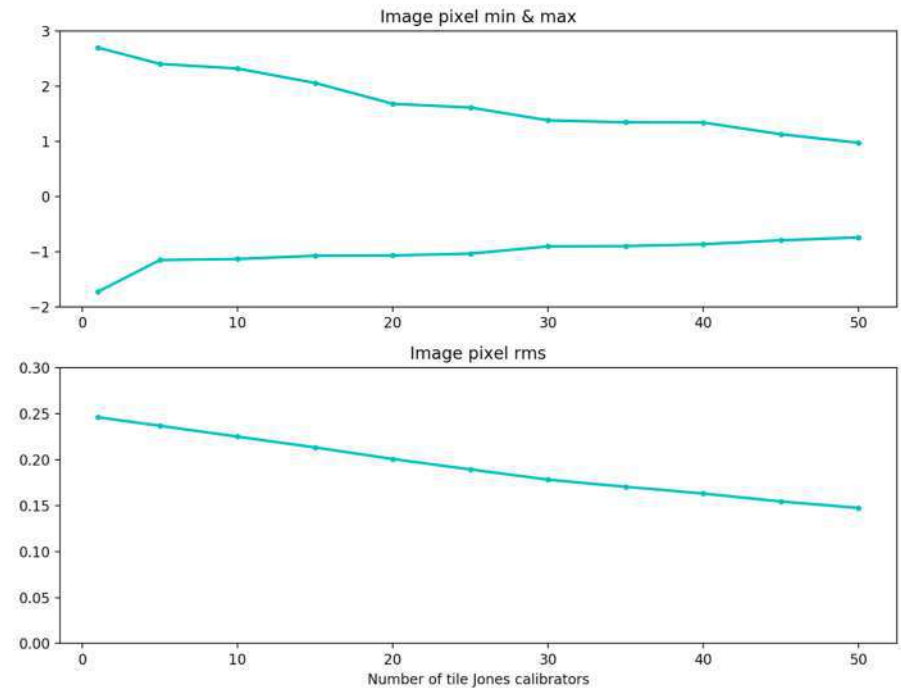

Peeling — over-peeling




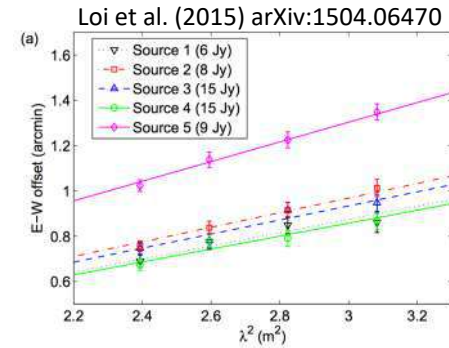

Dirty image with direction-independent calibration and 50 sources subtracted (all 50 peeled)

But as more of the 50 sources are peeled the weaker sources and their sidelobes disappear!


Peeling — over-peeling



Ionospheric Refraction



Constrained Peeling → linear phase model



Consider a set of visibilities with:


- a short time duration (≈ 10 sec)
- all-but calibrator c subtracted

$$V_{bf,c} \approx N_{bf} + M_{bf,c} \exp \left\{ -i2\pi\lambda_f^2 \left(u_{bf}\alpha_{l,c} + v_{bf}\alpha_{m,c} \right) \right\}$$

For small offsets ($\lambda^2\alpha \ll$ synthesised beam size) $\text{imag}\{V_{bf,c}\}$ is approximately linear in the α parameters \Rightarrow linear least-squares.


$$\begin{bmatrix} V_{b_{lf_1},c}^{\mathfrak{S}} - M_{b_{lf_1},c}^{\mathfrak{S}} \\ \vdots \\ V_{b_{Nf_M},c}^{\mathfrak{S}} - M_{b_{Nf_M},c}^{\mathfrak{S}} \end{bmatrix} \approx -2\pi \begin{bmatrix} \lambda_{f_1}^2 M_{b_{lf_1},c}^{\mathfrak{R}} u_{b_{lf_1}} & \lambda_{f_1}^2 M_{b_{lf_1},c}^{\mathfrak{R}} v_{b_{lf_1}} \\ \vdots & \vdots \\ \lambda_{f_M}^2 M_{b_{Nf_M},c}^{\mathfrak{R}} u_{b_{Nf_M}} & \lambda_{f_M}^2 M_{b_{Nf_M},c}^{\mathfrak{R}} v_{b_{Nf_M}} \end{bmatrix} \begin{bmatrix} \alpha_{l,c} \\ \alpha_{m,c} \end{bmatrix}$$

MWA data: 8 seconds \times ~ 1 MHz (182 MHz)



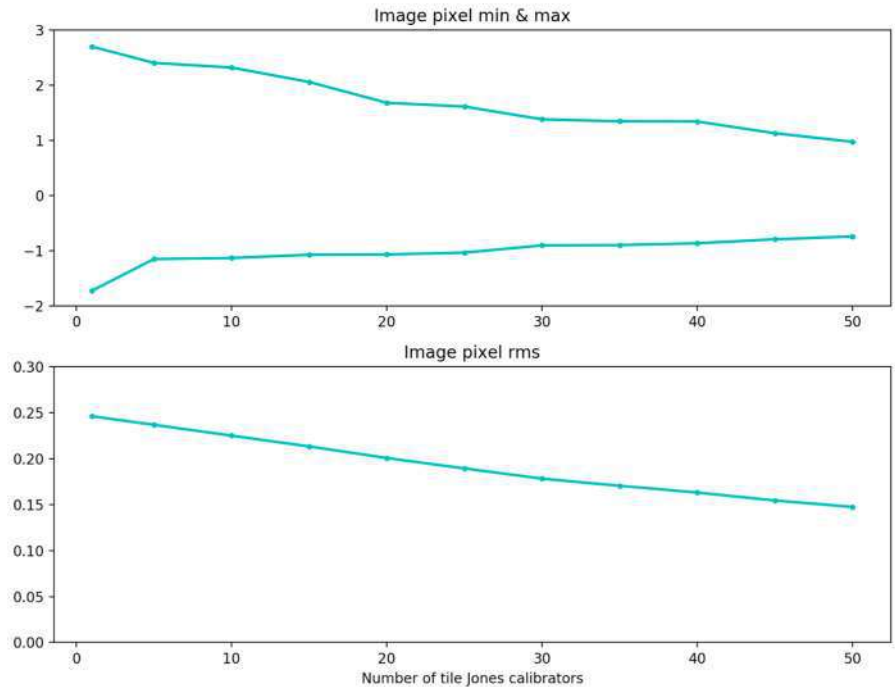

Colour scale ($[-1\sigma - +10\sigma]$):
[-0.19, 1.9] Jy/beam

MWA data: 8 seconds \times ~ 31 MHz (182 MHz)





Colour scale ($[-1\sigma - +10\sigma]$):
[-0.06, 0.6] Jy/beam


Peeling — over-peeling



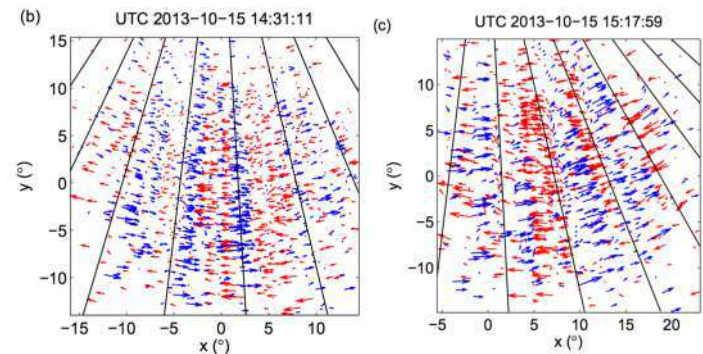

Peeling — over-peeling



Non-linear ionospheric phases




Cotton (2004) ASP Conf. Series 345, 74 MHz, 1-min VLA snapshots



Loi et al. (2015) arXiv:1504.06470


Constrained peeling → higher order models




$$\begin{bmatrix} V_{b_{f_1},c}^{\mathfrak{S}} - M_{b_{f_1},c}^{\mathfrak{S}} \\ \vdots \\ V_{b_{Nf_M},c}^{\mathfrak{S}} - M_{b_{Nf_M},c}^{\mathfrak{S}} \end{bmatrix} \approx -2\pi \begin{bmatrix} \lambda_{f_1}^2 M_{b_{f_1},c}^{\Re} u_{b_{f_1}} & \lambda_{f_1}^2 M_{b_{f_1},c}^{\Re} v_{b_{f_1}} \\ \vdots & \vdots \\ \lambda_{f_M}^2 M_{b_{Nf_M},c}^{\Re} u_{b_{Nf_M}} & \lambda_{f_M}^2 M_{b_{Nf_M},c}^{\Re} v_{b_{Nf_M}} \end{bmatrix} \begin{bmatrix} \alpha_{l,c} \\ \alpha_{m,c} \end{bmatrix}$$

Adding higher-order phase terms is straightforward:

$$V_{ijf,c}^{\mathfrak{S}} - M_{ijf,c}^{\mathfrak{S}} \approx -2\pi\lambda_f \left(dx_{ijf} \alpha_{l,c} + \left(x_{if}^2 - x_{jf}^2 \right) \beta_{l,c} + dy_{ijf} \alpha_{m,c} + \left(y_{if}^2 - y_{jf}^2 \right) \beta_{m,c} \right)$$

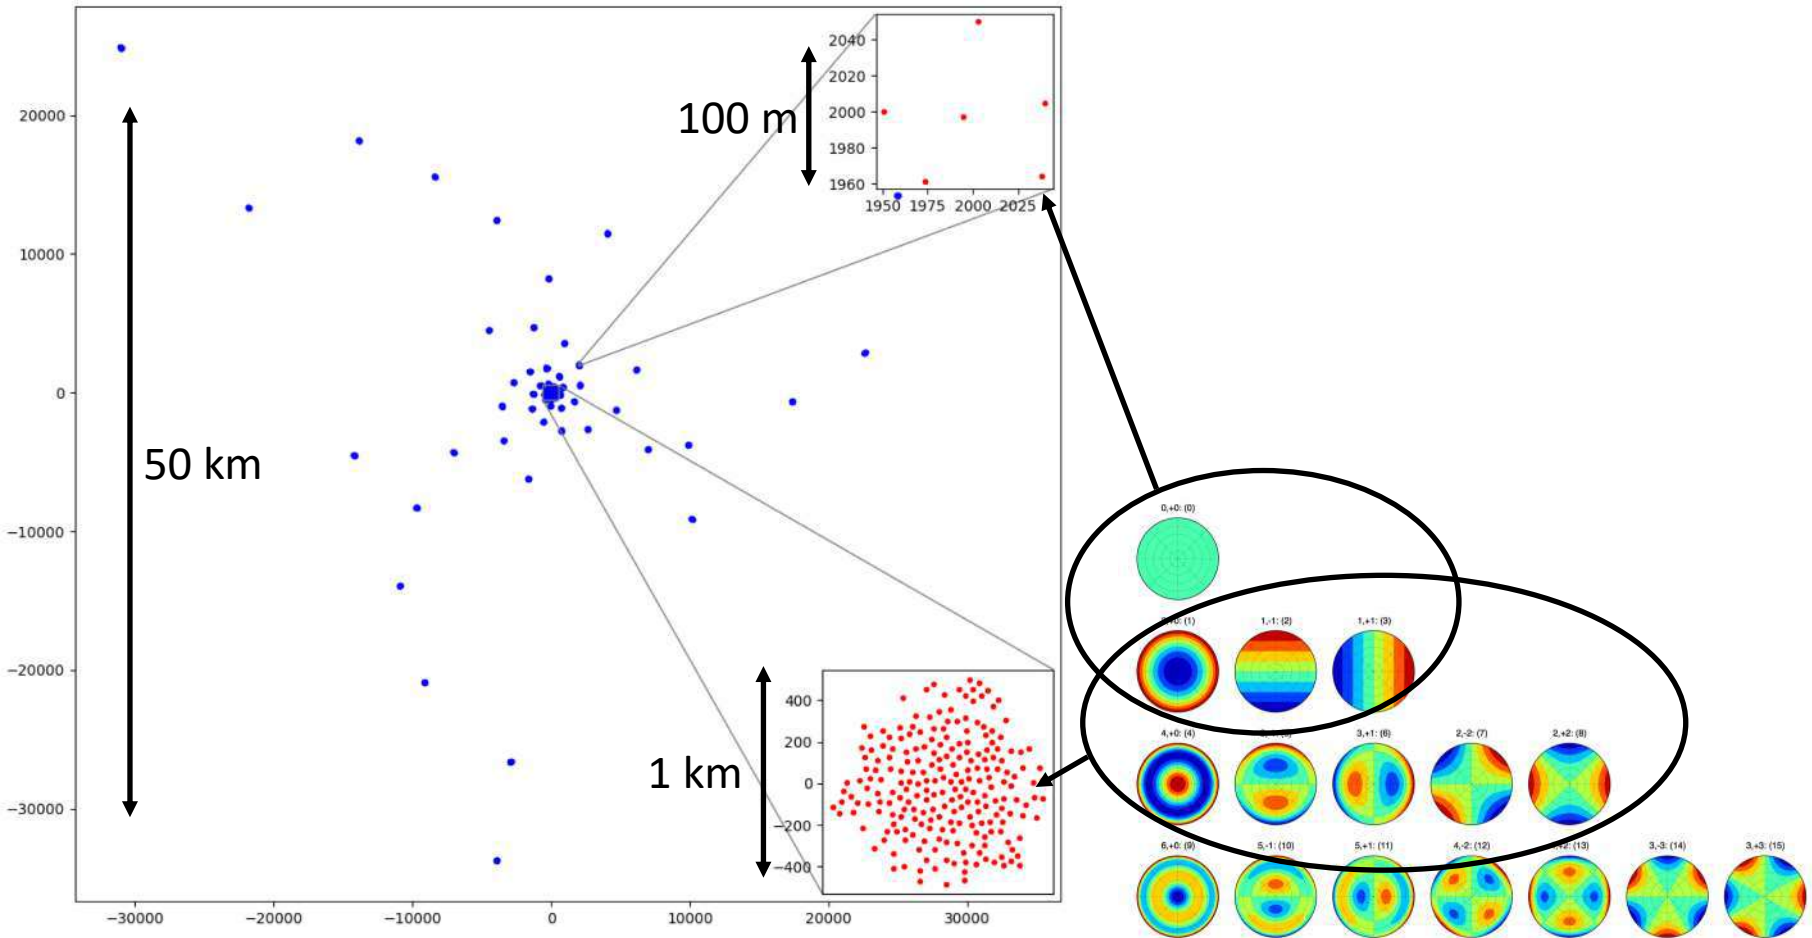


Or use polynomials




Constrained peeling → higher order models


SKA1-LOW Layout




Warped Snapshots



- Image warp is known and removed via image regridding
 - Potentially also the ionospheric perturbations.
- Time consuming to accurately calculate coordinates
 - Approximate: flat sky or interpolation
- Deep integrations occur in the image domain
 - With primary beam weighting, as in mosaicking




Simulated data: field centre: -3.5 to +3.5 hrs



Variable Polarised Primary Beams

All-sky polarised
tile response

Fourier
response



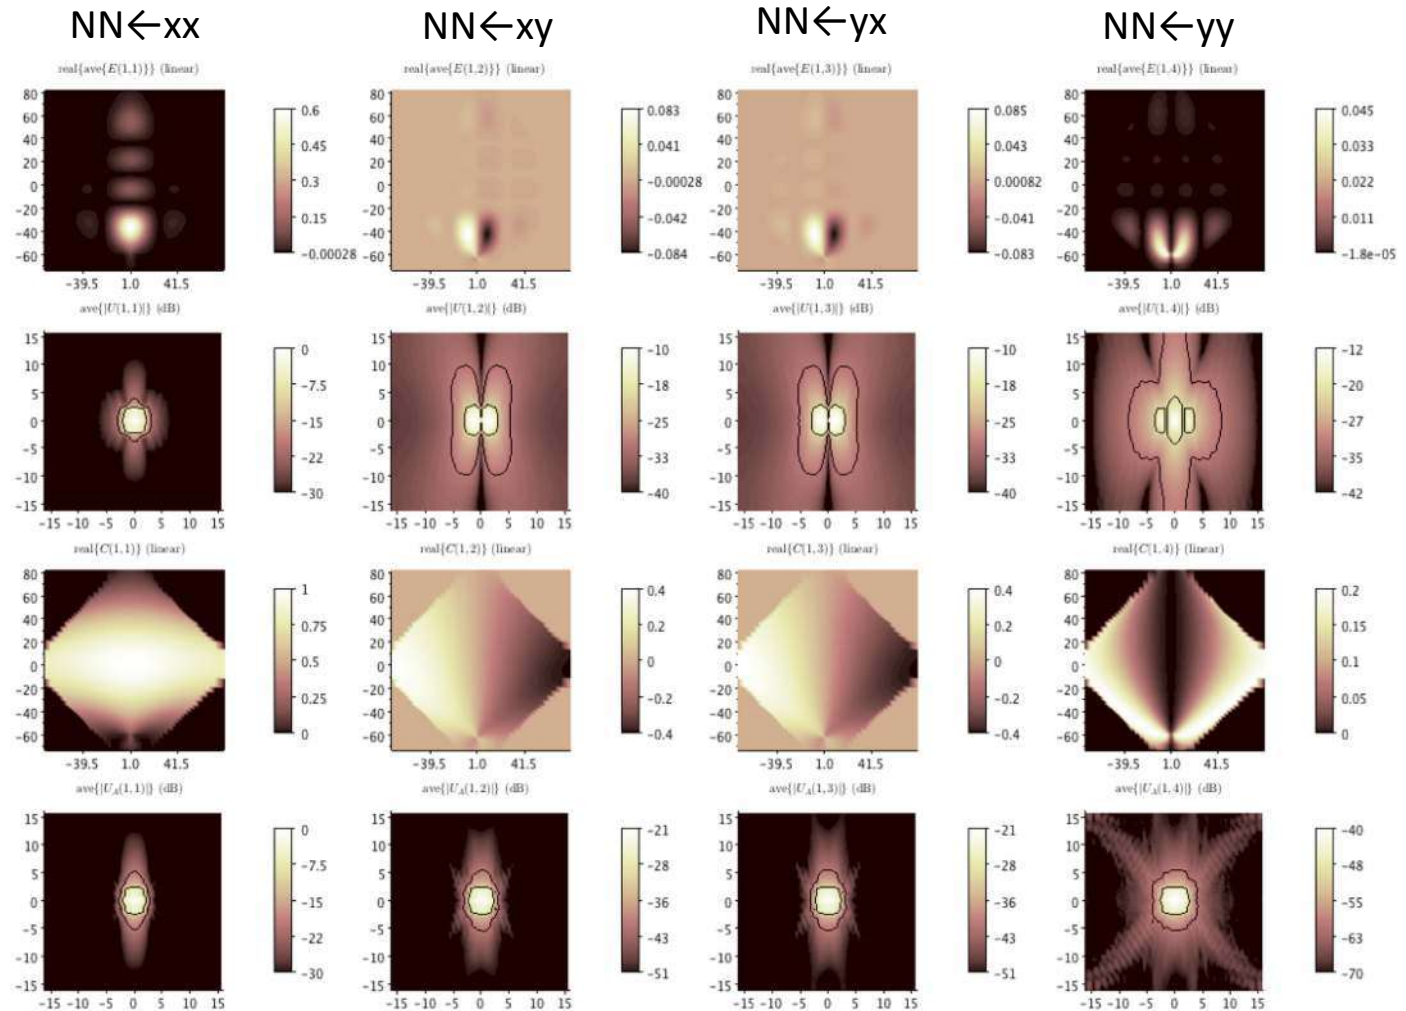
Variable Polarised Primary Beams

All-sky polarised
tile response

Fourier
response

Deal with curved sky
in the image domain


Remaining
Instrument Fourier
response



A projection


Apply direction-dependent corrections & weighting during gridding.

Not often used in practice, so still somewhat experimental.



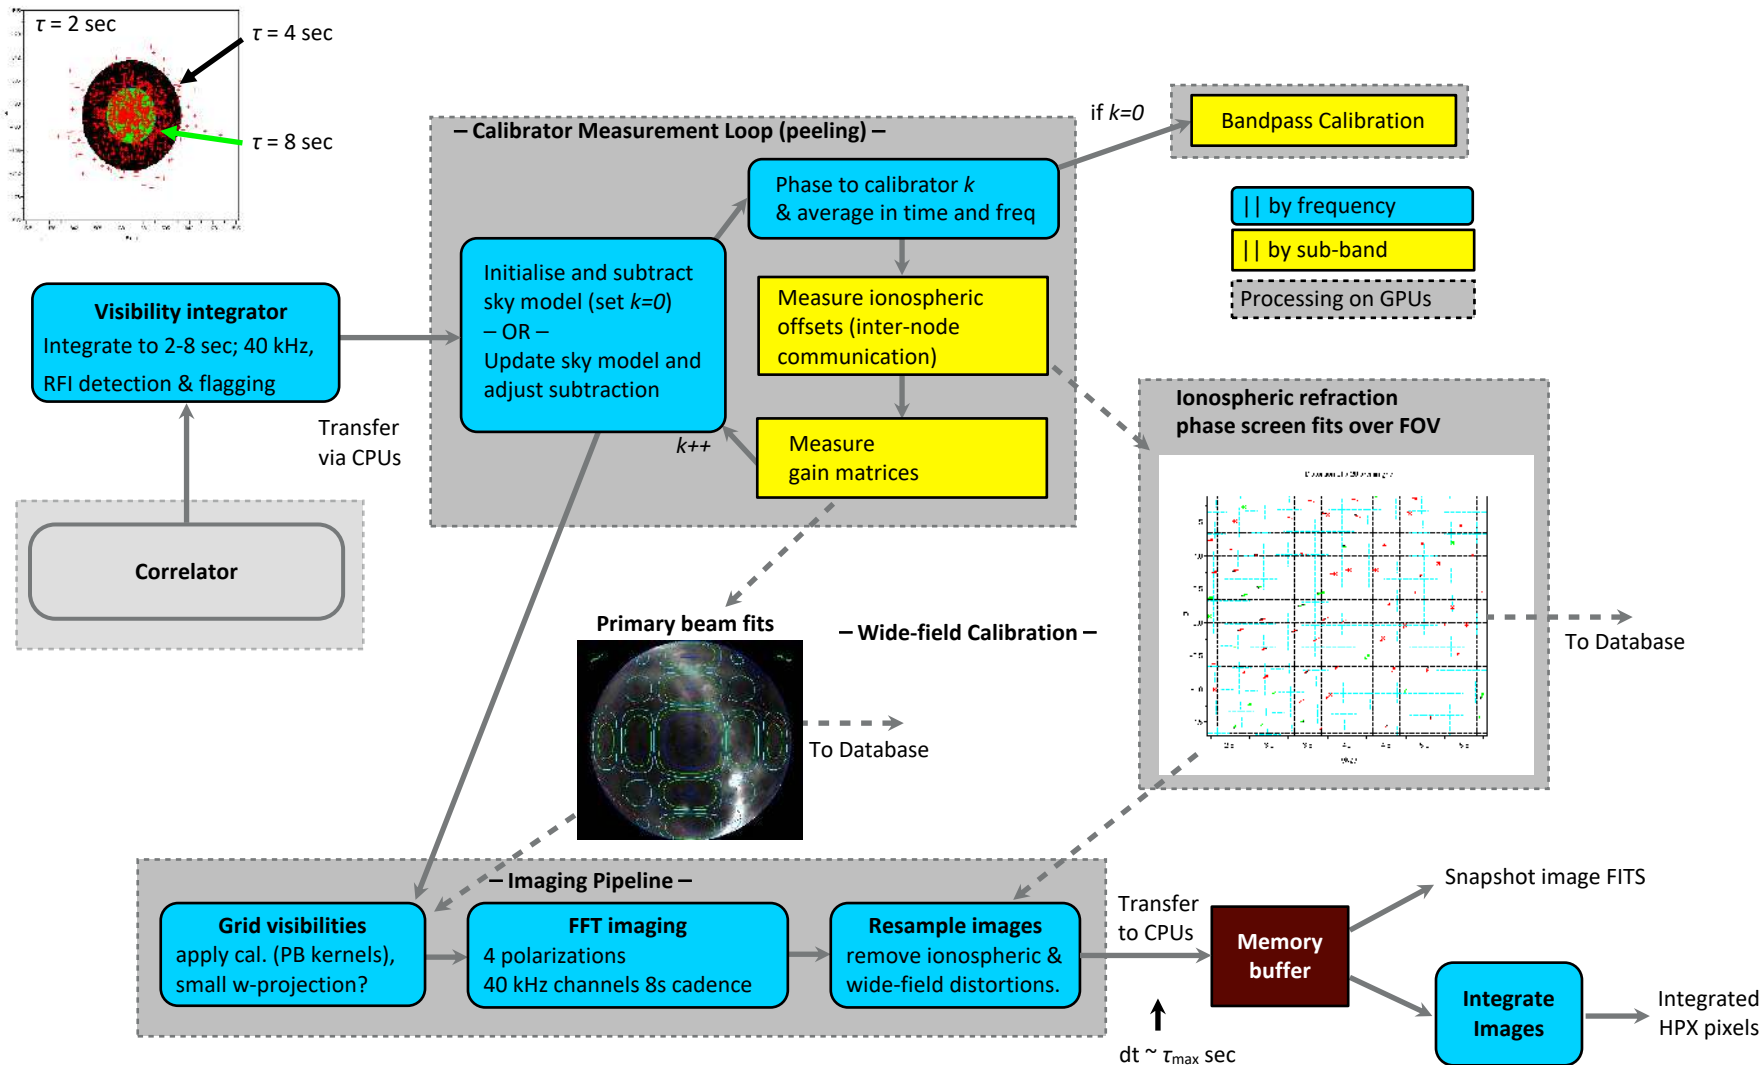
instrument pol \rightarrow instrument frame
 \Rightarrow compact in Fourier space

Response of one pol to unpolarised emission

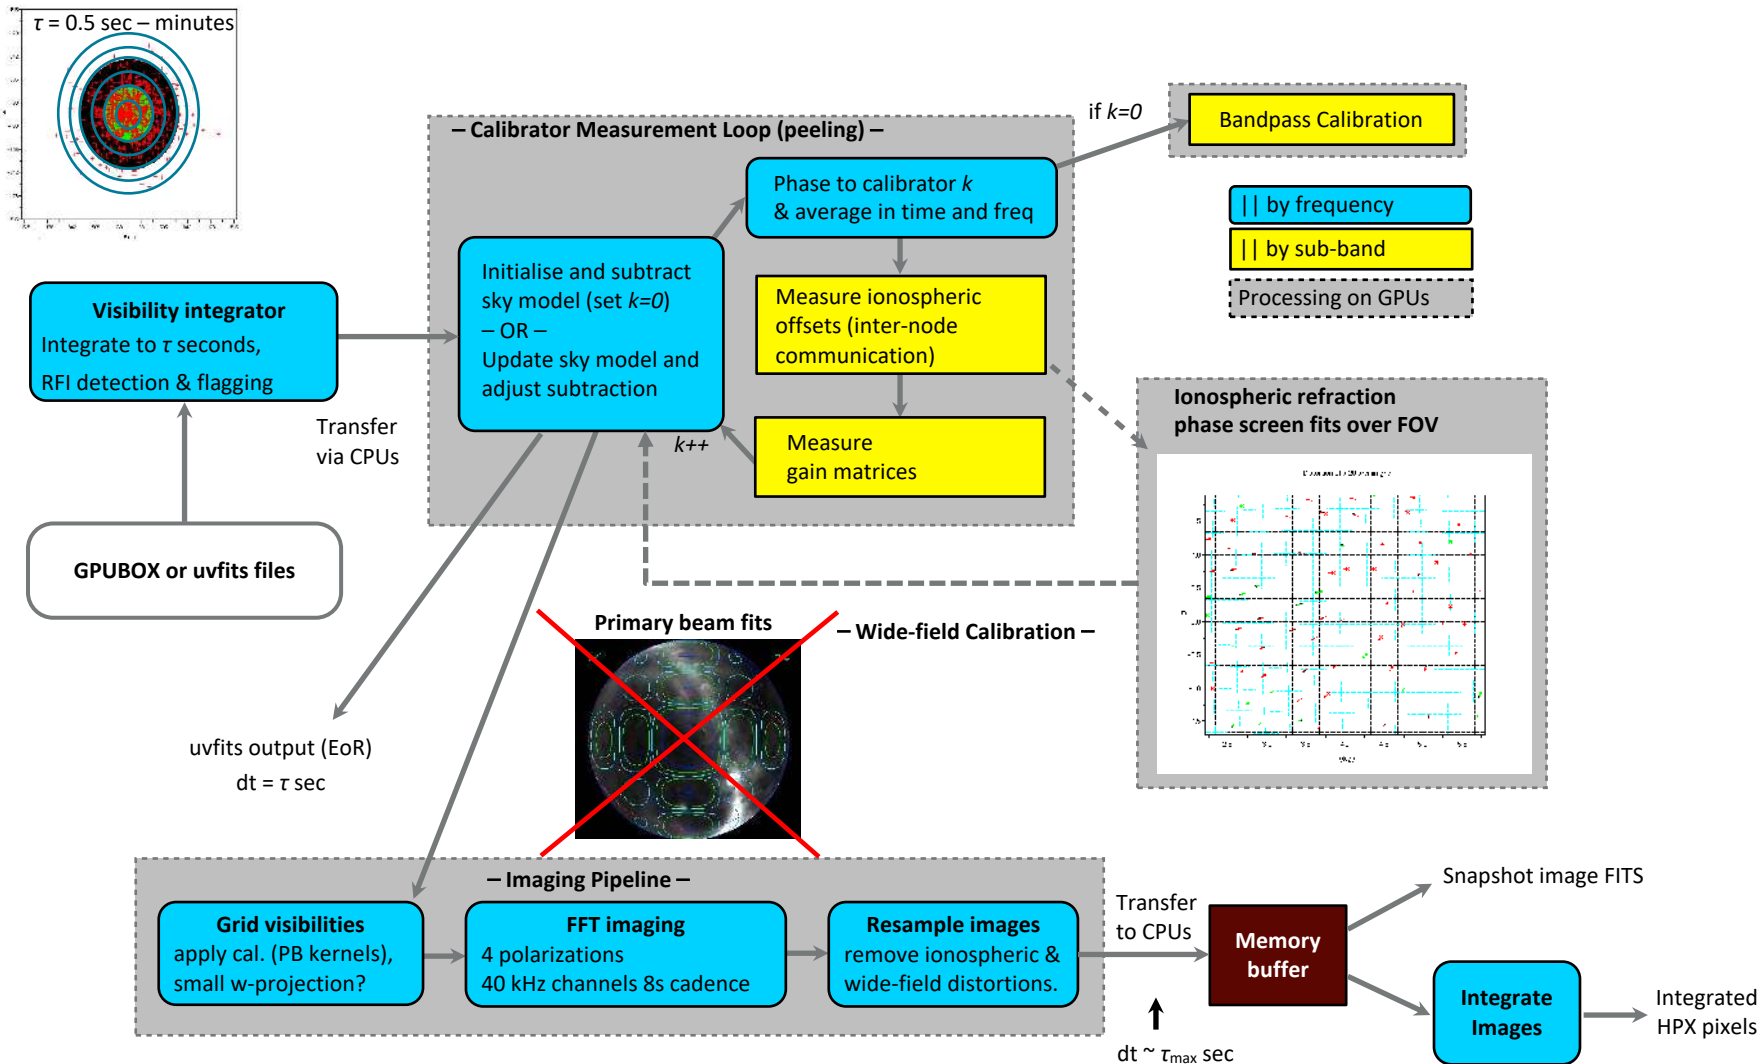


Data & Processing Flow

RTS Data Flow (original design)



RTS Data Flow (current design)




Heterogeneous Architecture



Galaxy@Pawsey 64 GPU nodes:

- Sandy Bridge EP CPUs
- K20X Kepler GPUs



Typical Processing Steps

- Process each 2 – 5 minute obsid separately
- First run
 - Set the solution interval, τ_{max} , to be the full obsid ($\approx 2 - 5$ minutes).
 - Set the sky model to have a dominant calibrator that contains all components in the field of view.
 - Generate bandpass calibration solutions.
- Frequency / obsid consensus? (continuous, smooth solutions ...)
- Second run
 - Set visibility integrator maximum to $\tau \approx 8$ seconds
 - Set sky model to have ≈ 1000 sources.
 - Subtract sky model with direction-dependent ionospheric calibration
 - Subtract a few sources with tile Jones matrix peeling if need be.
- Post-processing
 - Further image integration in time and/or frequency.
 - Conversion to Stokes images and FITS format.

Parameter Input Files

for obsid 1061313984

ImportCotterBasename=../1061313984/1061313984

MetafitsFilename=../1061313984/1061313984

CorrDumpTime=0.5

CorrDumpsPerCadence=128

NumberOfIntegrationBins=8

NumberOfIterations=1

BaseFilename=../1061313984/*_gpubox

ObservationTimeBase=2456528.22648


ObservationFrequencyBase=167.035

ChannelBandwidth=0.04

calBaselineMin=20.0

calShortBaselineTaper=40.0

First Run



% rts_node_gpu config_file.in

MakeImage=1

FieldOfViewDegrees=20

ImageOversampling=4


ObservationImageCentreRA=0.0

ObservationImageCentreDec=-27.0

FscrunchChan=32

Run a single RTS worker node.
(process a single 1.28 MHz coarse channel)
Use "rts_gpu" for full MPI version.

First Run



```
% python srclist_by_beam.py \  
-s srclist_puma-v2_complete.txt \  
-n 300 \  
-m ${obs}_metafits_ppds.fits
```


```
% rts_node_gpu config_file.in
```

```
generateDIjones=1  
useStoredCalibrationFiles=0
```

```
SourceCatalogueFile=patch300.txt
```

```
NumberOfCalibrators=1
```


First Run



```
% python srclist_by_beam.py \  
-s srclist_puma-v2_complete.txt \  
-n 300 \  
-m ${obs}_metafits_ppds.fits
```


```
% rts_node_gpu config_file.in
```

```
generateDIjones=1  
useStoredCalibrationFiles=0
```

```
SourceCatalogueFile=patch300.txt
```

```
NumberOfCalibrators=1  
NumberOfSourcesToPeel=1
```

First Run



```
% python srclist_by_beam.py \  
-s srclist_puma-v2_complete.txt \  
-n 300 \  
-m ${obs}_metafits_ppds.fits
```

```
% rts_node_gpu config_file.in
```


```
generateDIjones=1  
useStoredCalibrationFiles=0
```

```
SourceCatalogueFile=patch300.txt
```

```
NumberOfCalibrators=1  
NumberOfSourcesToPeel=0
```

```
imgBaselineMin=20.0  
imgShortBaselineTaper=40.0
```

First Run



```
% python srclist_by_beam.py \  
-s srclist_puma-v2_complete.txt \  
-n 300 \  
-m ${obs}_metafits_ppds.fits
```

```
% rts_node_gpu config_file.in
```


```
generateDIjones=1  
useStoredCalibrationFiles=0
```

```
SourceCatalogueFile=patch300.txt
```

```
NumberOfCalibrators=1  
NumberOfSourcesToPeel=1
```

```
imgBaselineMin=20.0  
imgShortBaselineTaper=40.0
```


Second Run



```
% python srclist_by_beam.py \  
-s srclist_puma-v2_complete.txt \  
-o experimental -x \  
-n 300 \  
-m ${obs}_metafits_ppds.fits
```

```
% rts_node_gpu config_file.in
```

```
generateDIjones=0  
useStoredCalibrationFiles=1
```

```
SourceCatalogueFile=peel300.txt
```

```
NumberOfSourcesToPrePeel=300  
NumberOfCalibrators=5  
NumberOfIonoCalibrators=300  
UpdateCalibratorAmplitudes=1  
NumberOfSourcesToPeel=300
```

Using the RTS

RTS Use

- EoR calibration and foreground (i.e. sky model) subtraction
- Calibration and imaging of Galactic polarisation
- Calibration and imaging for transient searches
- Calibration for pulsars beam-forming
- Calibration and imaging for other arrays (LEDA)

Running on Galaxy

slurm script:

```
#!/bin/bash -l
#SBATCH --job-name="RTS"
#SBATCH -o RTS-%A.out
#SBATCH --nodes=25
#SBATCH --ntasks-per-node=1
#SBATCH --time=00:20:00
#SBATCH --partition=gpuq
#SBATCH --account=mwaeor
#SBATCH --export=NONE
#SBATCH --mem=30000
#SBATCH --gres=gpu:1
```

module load rts

srun --ntasks=25 --ntasks-per-node=1 --export=ALL rts_gpu rts_params.in

Parameter Input Files

for obsid 1061313984

ImportCotterBasename=../1061313984/1061313984

MetafitsFilename=../1061313984/1061313984

CorrDumpTime=0.5

CorrDumpsPerCadence=128

NumberOfIntegrationBins=8

NumberOfIterations=1

BaseFilename=../1061313984/*_gpubox

ObservationTimeBase=2456528.22648

ObservationFrequencyBase=167.035

ChannelBandwidth=0.04

calBaselineMin=20.0

calShortBaselineTaper=40.0

Sky Catalogue Files

point source 1

SOURCE <name> ra_hrs dec_degs

FREQ freq_MHz I Q U V

FREQ freq_MHz I Q U V

...

ENDSOURCE

point source 2

SOURCE <name> ra_hrs dec_degs

FREQ freq_MHz I Q U V

...

ENDSOURCE

...


Sky Catalogue Files

Gaussian source

```
SOURCE <name> ra_hrs dec_degs  
GAUSSIAN PA_degs major_arcmin minor_arcmin  
FREQ freq_MHz I Q U V  
ENDSOURCE
```

Shapelet source

```
SOURCE <name> ra_hrs dec_degs  
FREQ freq_MHz I Q U V  
SHAPELET PA_degs major_arcmin minor_arcmin  
COEFF i1 j1 f1  
COEFF i2 j2 f2  
...  
COEFF iN jN fN  
ENDSOURCE
```



Fornax A
shapelet model

Sky Catalogue Files

Multi-component source

SOURCE <name> ra_hrs dec_degs
GAUSSIAN PA_degs major_arcmin minor_arcmin
FREQ freq_MHz I Q U V

COMPONENT ra_hrs dec_degs
FREQ freq_MHz I Q U V
ENDCOMPONENT

COMPONENT ra_hrs dec_degs
FREQ freq_MHz I Q U V
SHAPELET PA_degs major_arcmin minor_arcmin
COEFF i1 j1 f1
COEFF i2 j2 f2
...
COEFF iN jN fN
ENDCOMPONENT
ENDSOURCE

Flag Files

- Cotter flagging comes with the data.
- Extra flagging of tiles and/or frequency channels is available.
- Each MPI node will look for “flagged_tiles.txt” and “flagged_channels.txt” files and add extra flags.
- Very simple files:
 - single integer per line, representing to tile or channel to flag.
 - Integers start at zero and corresponds to input order
- Will be deprecate at some point, or advanced to contain metadata
 - But have been saying that for years, so mentioning here.

Summary

- RTS is very good at some things, and fast, but limited in scope
- To become a user, it is probably best to get in touch with me or one of the other groups using it:
 - EoR calibration and foreground (i.e. sky model) subtraction
 - Calibration and imaging of Galactic polarisation
 - Calibration and imaging for transient searches
 - Calibration for pulsars beam-forming
 - Calibration and imaging for other arrays (LEDA)